

Solution techniques in OpenFOAM

Artur Lidtke

University of Southampton
akl1g09@soton.ac.uk

November 4, 2014

Content

Concepts

- Solution ideology
- Understanding the process
- Available methods

Practice - flat plate

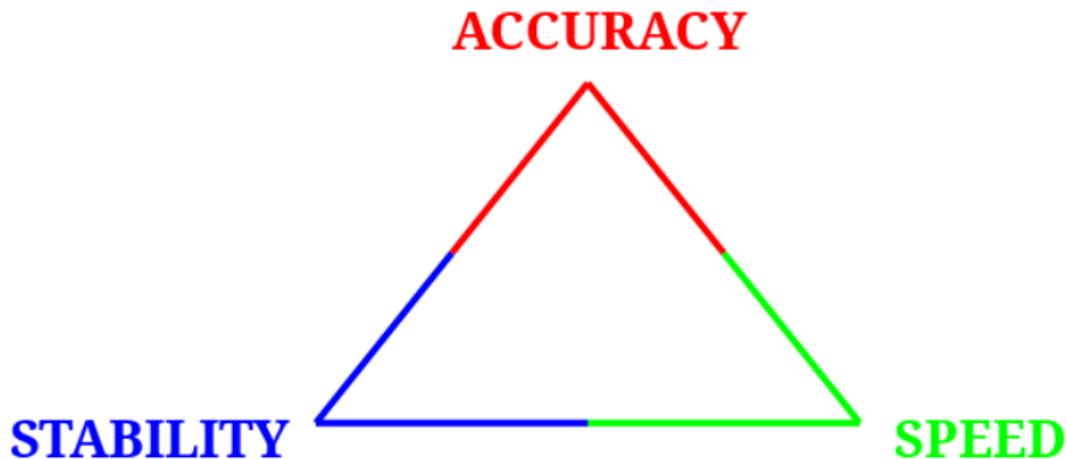
- Test case overview
- Refining the mesh
- Under-relaxation
- High order schemes

Summary



Concepts

The balance



Choosing the correct set-up

- ▶ Cannot have maximum accuracy with rock-solid stability at a minimum cost
 - ▶ Need to trade off the three drivers based on the application
 - ▶ There is no definite answer that will always work for all problems, even vaguely similar ones
 - ▶ Check forums, read papers and look at tutorials to find a starting point
 - ▶ Then run a sensitivity study to see what works best for you
 - ▶ Often useful to use a different set-up at various stages of a simulation - high stability first for an initial estimate and then move towards accuracy
 - ▶ Stay up to date with code development to see what's new

How does this work?

- ▶ We use a set of small, finite volumes to describe our domain of interest
- ▶ Numerical **schemes** of various orders are used to construct linear equations approximating our problem
- ▶ The resulting matrices are then given to linear **solvers** which try to satisfy the equations
- ▶ An indication of how well an equation is satisfied is its **residual**
- ▶ This is done for all equations in turn in order to resolve the **pressure-velocity coupling**
- ▶ Need to repeat this at each time step (unsteady) or iteration (steady)

What's the point of all this stuff in the logs? - SIMPLE

```

smoothSolver: Solving for Ux, Initial residual = 9.69841e-08, Final residual = 9.69841e-08, No Iterations 0
smoothSolver: Solving for Uy, Initial residual = 7.00898e-07, Final residual = 7.94086e-08, No Iterations 3

GAMG: Solving for p, Initial residual = 0.0032684, Final residual = 2.99124e-06, No Iterations 26
GAMG: Solving for p, Initial residual = 3.00459e-06, Final residual = 9.44095e-08, No Iterations 60
GAMG: Solving for p, Initial residual = 9.44092e-08, Final residual = 9.44092e-08, No Iterations 0
time step continuity errors : sum local = 2.32813e-10, global = 3.76108e-12, cumulative = 8.16684e-12

smoothSolver: Solving for omega, Initial residual = 9.81049e-08, Final residual = 9.81049e-08, No Iterations 0
smoothSolver: Solving for k, Initial residual = 1.21105e-07, Final residual = 6.45996e-08, No Iterations 1

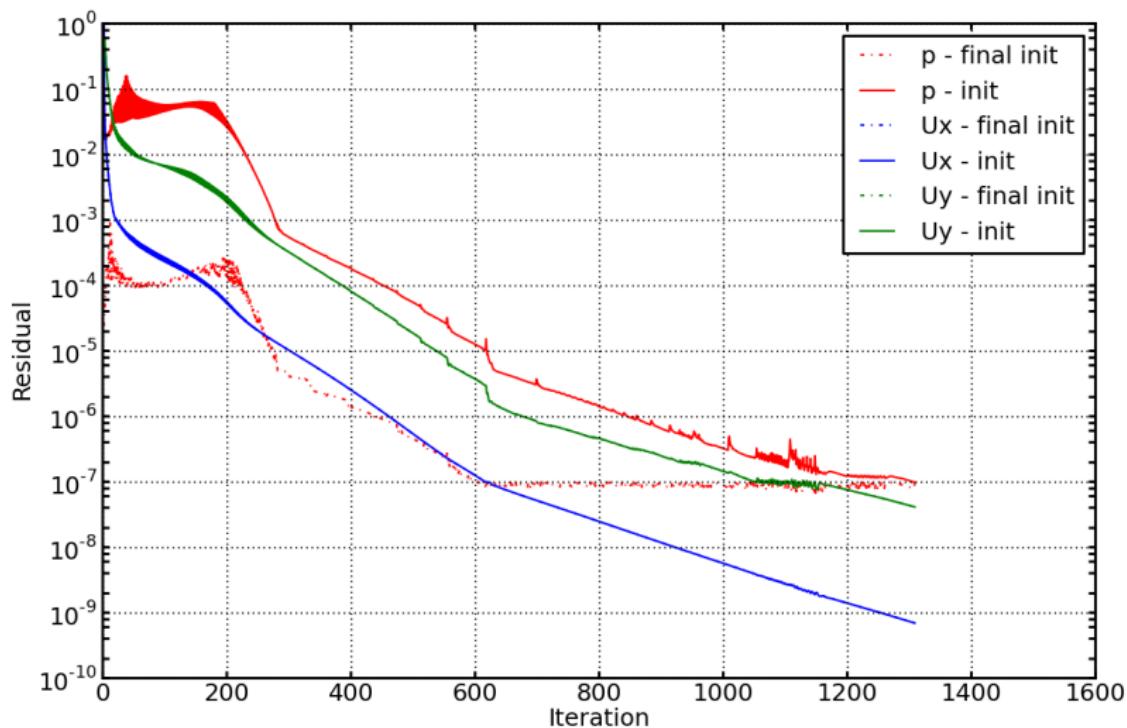
```

Steady-state solution (SIMPLE)

- ▶ Start with no information about the flow (maybe a potential solution)
 - ▶ Iterate over velocity and pressure interchangeably to drive the equations towards convergence
 - ▶ Need to monitor the residuals of all equations to make sure the approximate solution satisfies them
 - ▶ May control **underrelaxation** in order to "dampen" the solution process - for $\alpha < 1$ more stable but slower convergence

$$\phi^n = \phi^{n-1} + \alpha(\phi^{n*} - \phi^{n-1})$$

Graphically speaking



What's the point of all this stuff in the logs? - PISO

Courant Number mean: 0.0976805 max: 0.585722

DILUPBiCG: Solving for Ux, Initial residual = 0.148584,
Final residual = 7.15711e-06, No Iterations 6

DILUPBiCG: Solving for Uy, Initial residual = 0.256618,
Final residual = 8.94127e-06, No Iterations 6

DICPCG: Solving for p, Initial residual = 0.379232, Final
residual = 3.38648e-07, No Iterations 34

time step continuity errors : sum local = 3.15698e-09,
global = 1.94222e-19, cumulative = 9.28841e-19

DICPCG: Solving for p, Initial residual = 0.286937, Final
residual = 5.99637e-07, No Iterations 33

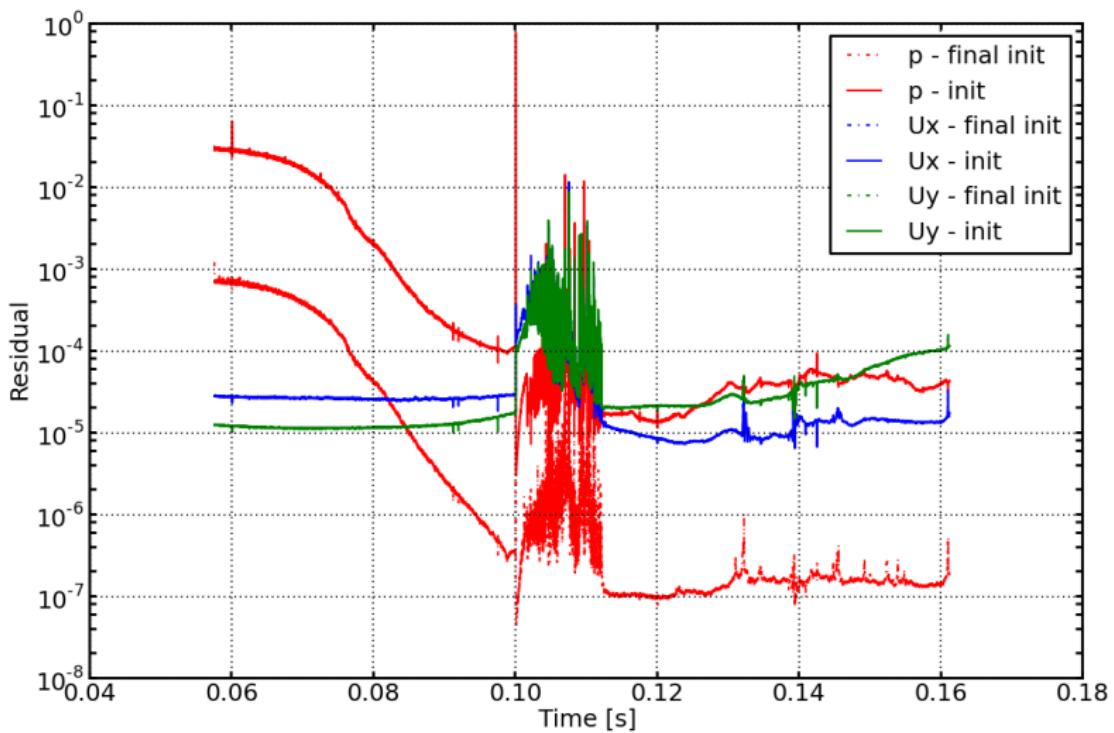
time step continuity errors : sum local = 6.08774e-09,
global = 5.80474e-19, cumulative = 1.50932e-18

Unsteady solution (PISO/PIMPLE)

- ▶ Need to reduce the residuals **at each time step** and not just in the long run as in SIMPLE
- ▶ May be very hard to maintain stability without a good initial condition
- ▶ Must monitor the entire solution, not just the residuals, to ensure convergence at each time step
- ▶ Check if the solvers reach maximum no. iterations before the residuals drop by the required amount
- ▶ If yes then
 - ▶ Reduce the time step
 - ▶ Increase the no. solver iterations
 - ▶ Use more loops over each equation (`nCorrectors`, `nOuterCorrectors`, etc. in `fvSolution`)

Understanding the process

Graphically speaking



Courant number

- ▶ At each time the Courant number is computed for an unsteady problem

Courant number

- ▶ At each time the Courant number is computed for an unsteady problem

$$Co = \frac{u\Delta t}{\Delta x}$$

Δt - simulation time step

Δx - grid size

u - local flow velocity

OR

"How many cells does a fluid particle travel through in one time step"

Courant number

- ▶ At each time the Courant number is computed for an unsteady problem
- ▶ Should keep Co reasonable - depending on the schemes and methods there may be a theoretical limit for stability
- ▶ Generally, $Co < 1.0$ is a good rule of thumb, **will** need less for LES with high order schemes

Courant number

- ▶ At each time the Courant number is computed for an unsteady problem
- ▶ Should keep Co reasonable - depending on the schemes and methods there may be a theoretical limit for stability
- ▶ Generally, $Co < 1.0$ is a good rule of thumb, **will** need less for LES with high order schemes
- ▶ The higher the Reynolds number (u), the smaller the cells have to be to resolve the boundary layer

Courant number

- ▶ At each time the Courant number is computed for an unsteady problem
- ▶ Should keep Co reasonable - depending on the schemes and methods there may be a theoretical limit for stability
- ▶ Generally, $Co < 1.0$ is a good rule of thumb, **will** need less for LES with high order schemes
- ▶ The higher the Reynolds number (u), the smaller the cells have to be to resolve the boundary layer
- ▶ The smaller the cells and the higher the speed the smaller the time step that allows a stable solution

Courant number

- ▶ At each time the Courant number is computed for an unsteady problem
- ▶ Should keep Co reasonable - depending on the schemes and methods there may be a theoretical limit for stability
- ▶ Generally, $Co < 1.0$ is a good rule of thumb, **will** need less for LES with high order schemes
- ▶ The higher the Reynolds number (u), the smaller the cells have to be to resolve the boundary layer
- ▶ The smaller the cells and the higher the speed the smaller the time step that allows a stable solution
- ▶ **CFD of High-Re flows is still a huge challenge!**

Schemes - convective

- ▶ Full list of available options on
<http://www.openfoam.org/docs/user/fvSchemes.php>
- ▶ If you are going to use upwind for LES then don't bother!
- ▶ Typical choices for RANS/URANS
 - ▶ upwind - stable but inaccurate
 - ▶ linearUpwind - mixed 1st-2nd order, more accurate but less stable
- ▶ My suggestions for LES
 - ▶ linear - 2nd order accuracy but very unstable (convective term)
 - ▶ filteredLinear/limitedLinear - special mixed schemes supposed to balance accuracy and stability (convective term)

Schemes - time

- ▶ Two main choices:
 - ▶ Euler - 1st order time scheme, stable but not accurate - not a problem since URANS introduces even bigger inaccuracies!
 - ▶ backward/Crank-Nicholson - 2nd order time schemes - a must for LES
- ▶ The time scheme will influence the maximum time step greatly
- ▶ Typical symptom for an LES case running at a too large time step for a high order time scheme is the last loop over pressure equation exploding (residuals > 1, pressure field becoming unphysical)

Other schemes

- ▶ These include
 - ▶ Gradient
 - ▶ Laplacian
 - ▶ Surface normal gradient
 - ▶ Interpolation
- ▶ Generally affect stability less than the time and convective scheme
- ▶ As a rule of thumb OK to use linear unless a specific problem requires otherwise - look at literature and tutorial cases to get an idea

Other schemes

- ▶ These include
 - ▶ Gradient
 - ▶ Laplacian
 - ▶ Surface normal gradient
 - ▶ Interpolation
- ▶ Generally affect stability less than the time and convective scheme
- ▶ As a rule of thumb OK to use linear unless a specific problem requires otherwise - look at literature and tutorial cases to get an idea
- ▶ Let us examine the flat plate test case with the default settings I specified

Other schemes

- ▶ These include
 - ▶ Gradient
 - ▶ Laplacian
 - ▶ Surface normal gradient
 - ▶ Interpolation
- ▶ Generally affect stability less than the time and convective scheme
- ▶ As a rule of thumb OK to use linear unless a specific problem requires otherwise - look at literature and tutorial cases to get an idea
- ▶ Let us examine the flat plate test case with the default settings I specified

Go to

```
cd $FOAM_RUN/flatPlate/ZPG_wallFunction_coarse
```

Other schemes

- ▶ These include
 - ▶ Gradient
 - ▶ Laplacian
 - ▶ Surface normal gradient
 - ▶ Interpolation
- ▶ Generally affect stability less than the time and convective scheme
- ▶ As a rule of thumb OK to use linear unless a specific problem requires otherwise - look at literature and tutorial cases to get an idea
- ▶ Let us examine the flat plate test case with the default settings I specified

Edit

gedit system/fvSchemes

Solvers

- ▶ Even less obvious choice than the schemes, here are some commonly used ones:
 - ▶ smoothSolver - fairly stable, good for steady RANS (U , k , ω)
 - ▶ GAMG + GaussSeidel - robust for the pressure equation
 - ▶ PBiCG + DILU - works OK for velocity for unsteady problems, can also use for scalar fields (like volume fraction, for instance)
- ▶ Factors affecting the optimum solver for a given problem:
 - ▶ Size and nature of the mesh (structured/unstructured)
 - ▶ Parallelisation strategy
 - ▶ Numerical schemes used
 - ▶ Solution strategy (no. loops per equation etc.)
- ▶ More info on
<http://www.openfoam.org/docs/user/fvSolution.php>

Solvers

- ▶ Even less obvious choice than the schemes, here are some commonly used ones:
 - ▶ smoothSolver - fairly stable, good for steady RANS (U , k , ω)
 - ▶ GAMG + GaussSeidel - robust for the pressure equation
 - ▶ PBiCG + DILU - works OK for velocity for unsteady problems, can also use for scalar fields (like volume fraction, for instance)
- ▶ Factors affecting the optimum solver for a given problem:
 - ▶ Size and nature of the mesh (structured/unstructured)
 - ▶ Parallelisation strategy
 - ▶ Numerical schemes used
 - ▶ Solution strategy (no. loops per equation etc.)
- ▶ More info on
<http://www.openfoam.org/docs/user/fvSolution.php>

Edit

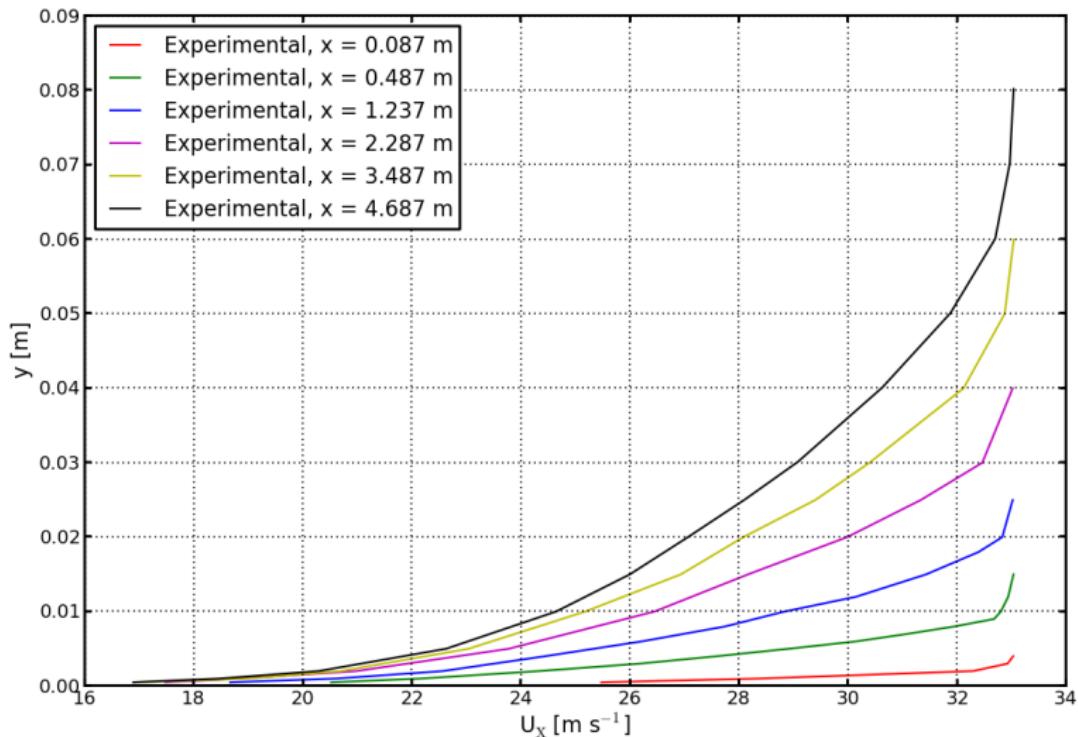
gedit system/fvSolution

Practice - flat plate

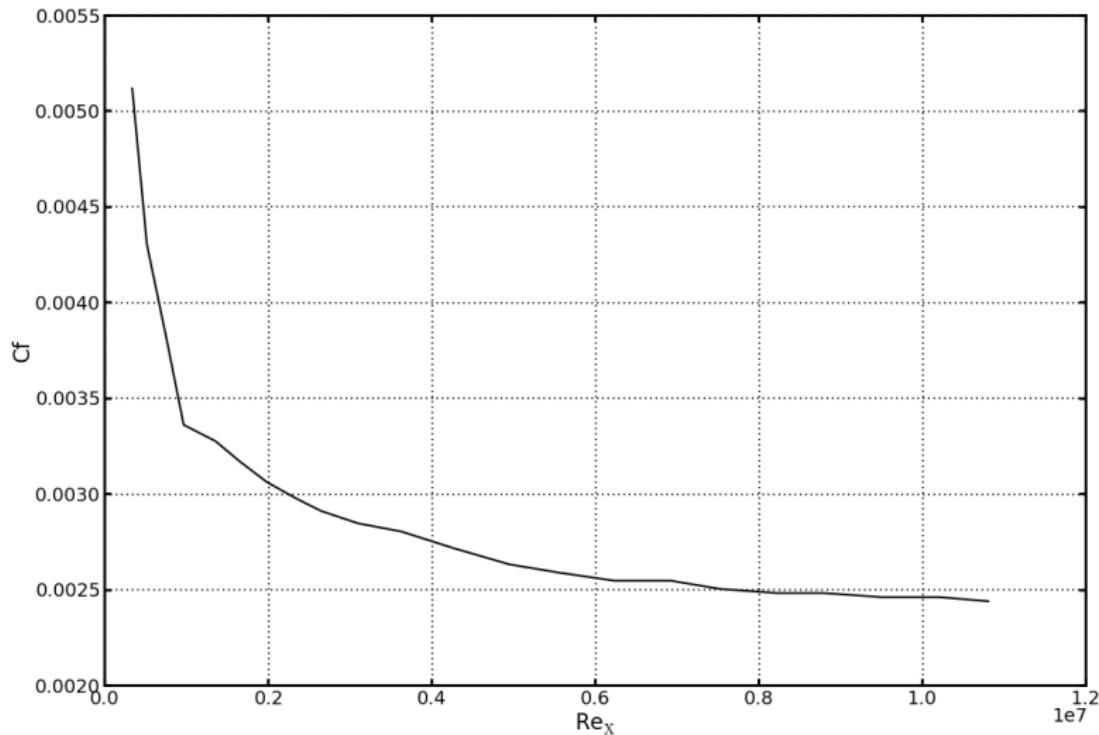
Test case

- ▶ Flat plate in zero pressure gradient condition
- ▶ Probably the most fundamental hydrodynamics problem
- ▶ Very relevant to all engineering applications
- ▶ Extensive amount of validation data in the public domain
 - ▶ Wieghardt et al. 1968
<http://www.grc.nasa.gov/WWW/wind/valid/fpturb/fpturb01/fpturb01.html>
 - ▶ T3A transitional Re regime; see for instance Cutrone et al. 2007
http://cfd.mace.manchester.ac.uk/twiki/pub/Main/ClareTurner/cutrone_bypass.pdf

Test case



Test case



Baseline simulation

- ▶ Let us examine the flat plate test case with the default settings I specified

Baseline simulation

- ▶ Let us examine the flat plate test case with the default settings I specified

Go to

```
cd $FOAM_RUN/flatPlate/ZPG_wallFunction_coarse
```

Baseline simulation

- ▶ Let us examine the flat plate test case with the default settings I specified
- ▶ There is a ready script that should execute the solution, let's have a look

Baseline simulation

- ▶ Let us examine the flat plate test case with the default settings I specified
- ▶ There is a ready script that should execute the solution, let's have a look



Baseline simulation

- ▶ Let us examine the flat plate test case with the default settings I specified
- ▶ There is a ready script that should execute the solution, let's have a look
- ▶ Now, run the script. The case should converge quickly

Baseline simulation

- ▶ Let us examine the flat plate test case with the default settings I specified
- ▶ There is a ready script that should execute the solution, let's have a look
- ▶ Now, run the script. The case should converge quickly

Use
./Allrun

Baseline simulation

- ▶ Let us examine the flat plate test case with the default settings I specified
- ▶ There is a ready script that should execute the solution, let's have a look
- ▶ Now, run the script. The case should converge quickly
- ▶ Let us take a look at the residuals and forces convergence

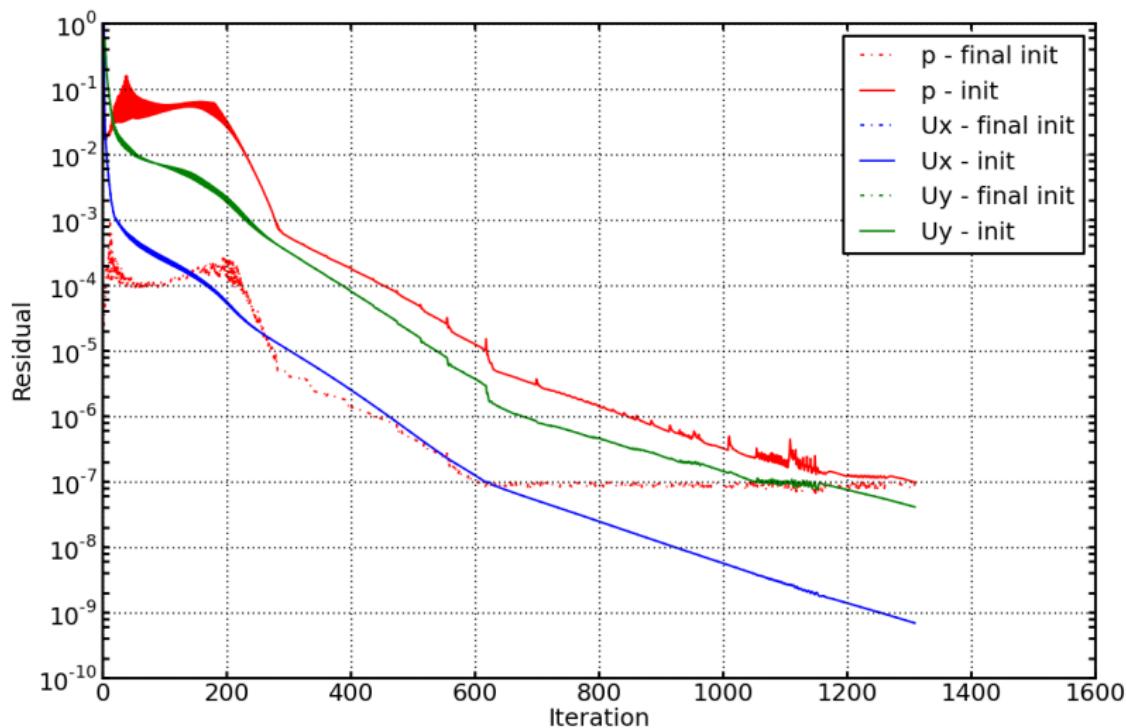
Baseline simulation

- ▶ Let us examine the flat plate test case with the default settings I specified
- ▶ There is a ready script that should execute the solution, let's have a look
- ▶ Now, run the script. The case should converge quickly
- ▶ Let us take a look at the residuals and forces convergence

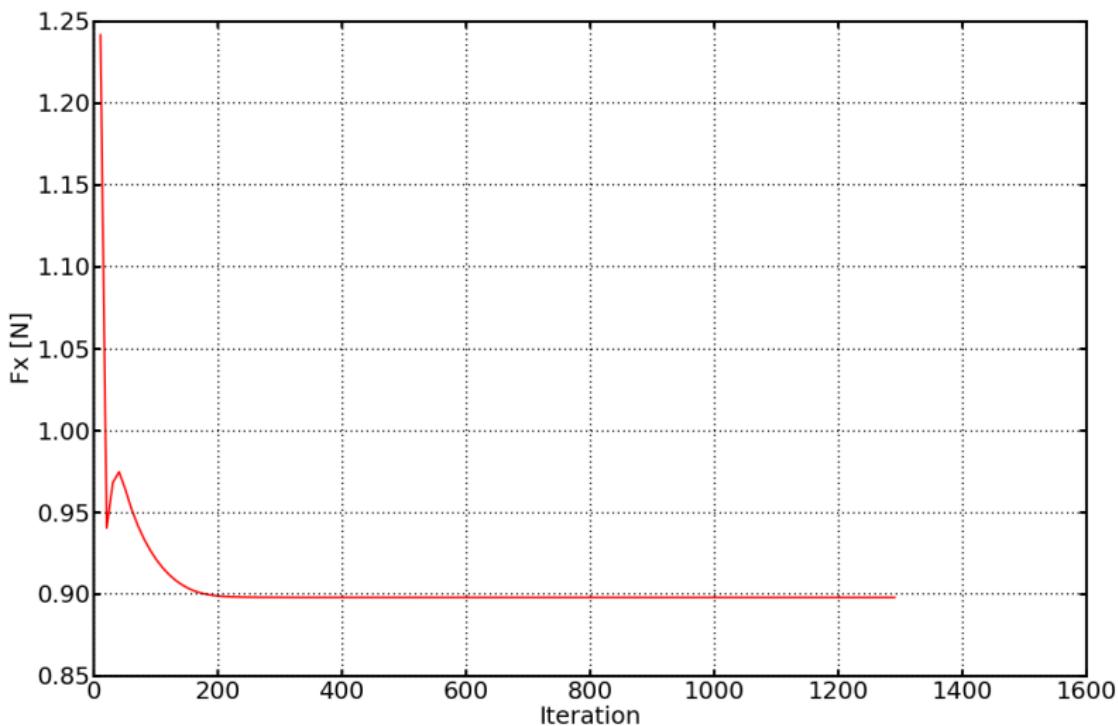
Use

- ▶ `cd ..`
- ▶ `python plotResiduals.py`
- ▶ `python plotForces.py`

Baseline simulation - residuals



Baseline simulation - drag force



What does this tell us?

- ▶ See both the pressure and velocity equations steadily progressing towards convergence
- ▶ Drag force settles quite quickly - the wall function makes convergence of the flow on the boundary easy but not necessarily very accurate!
- ▶ The remaining time of the simulation is spent on sorting out flow features
- ▶ Homework: we also collected boundary layer profile data, have a look at how these change as the solution progresses

Refining the mesh

- ▶ Let us use the coarse mesh solution to initialise one for a refined grid

Refining the mesh

- ▶ Let us use the coarse mesh solution to initialise one for a refined grid

Go to

```
cd ZPG_wallFunction_fine
```

Refining the mesh

- ▶ Let us use the coarse mesh solution to initialise one for a refined grid
- ▶ Note how we use the `mapFields` utility in the `Allrun` script to map the final fields from the coarse mesh onto the fine one

Refining the mesh

- ▶ Let us use the coarse mesh solution to initialise one for a refined grid
- ▶ Note how we use the `mapFields` utility in the `Allrun` script to map the final fields from the coarse mesh onto the fine one

Note

```
mapFields -sourceTime 2000 -consistent  
.../ZPG_wallFunction_coarse
```

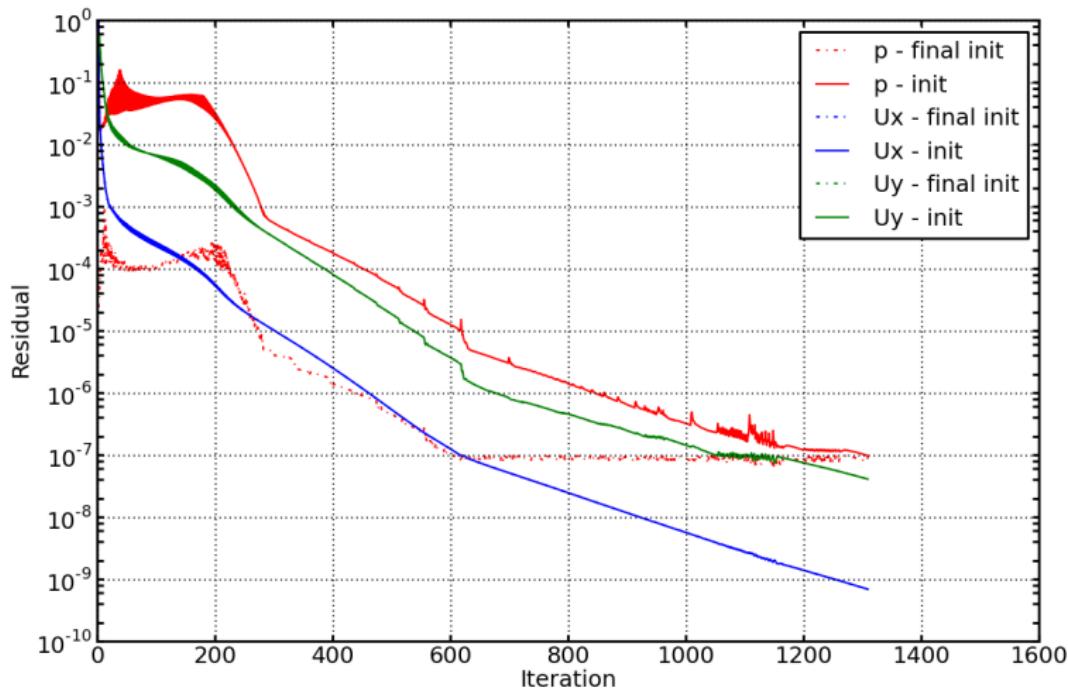
Refining the mesh

- ▶ Let us use the coarse mesh solution to initialise one for a refined grid
- ▶ Note how we use the `mapFields` utility in the `Allrun` script to map the final fields from the coarse mesh onto the fine one

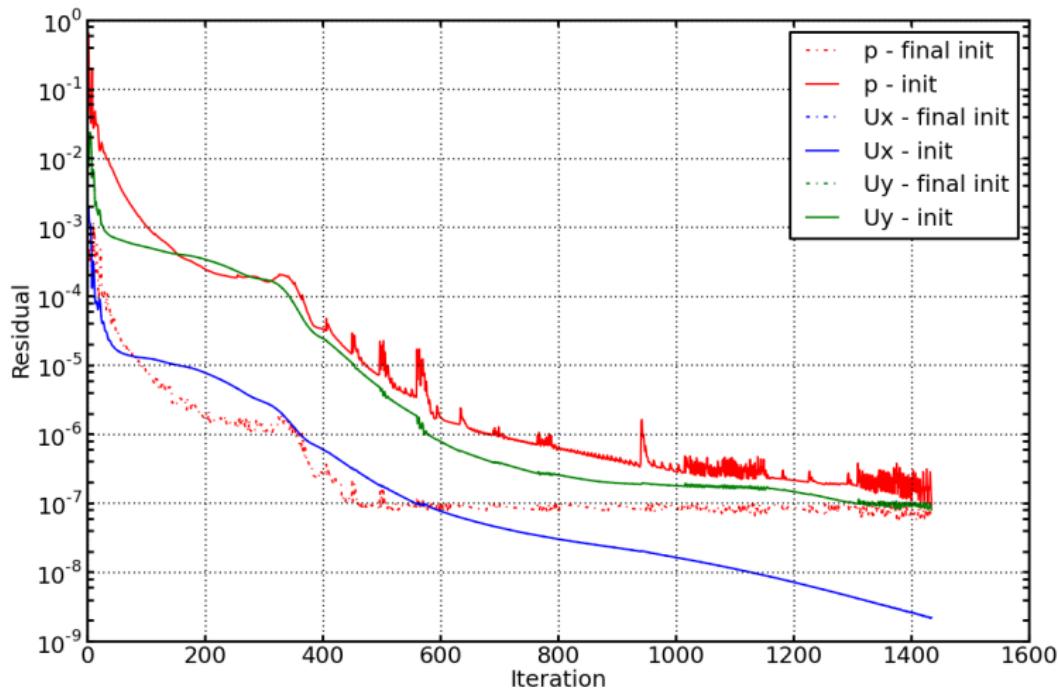
Do

- ▶ Run the simulation
- ▶ Look at the forces and residuals

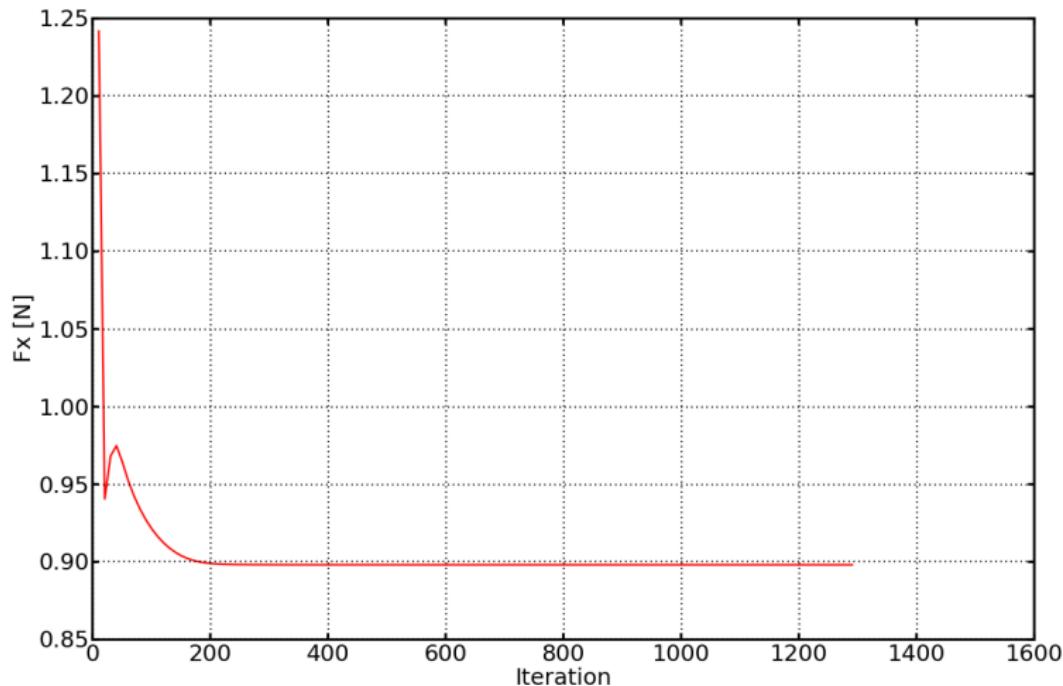
Original - residuals



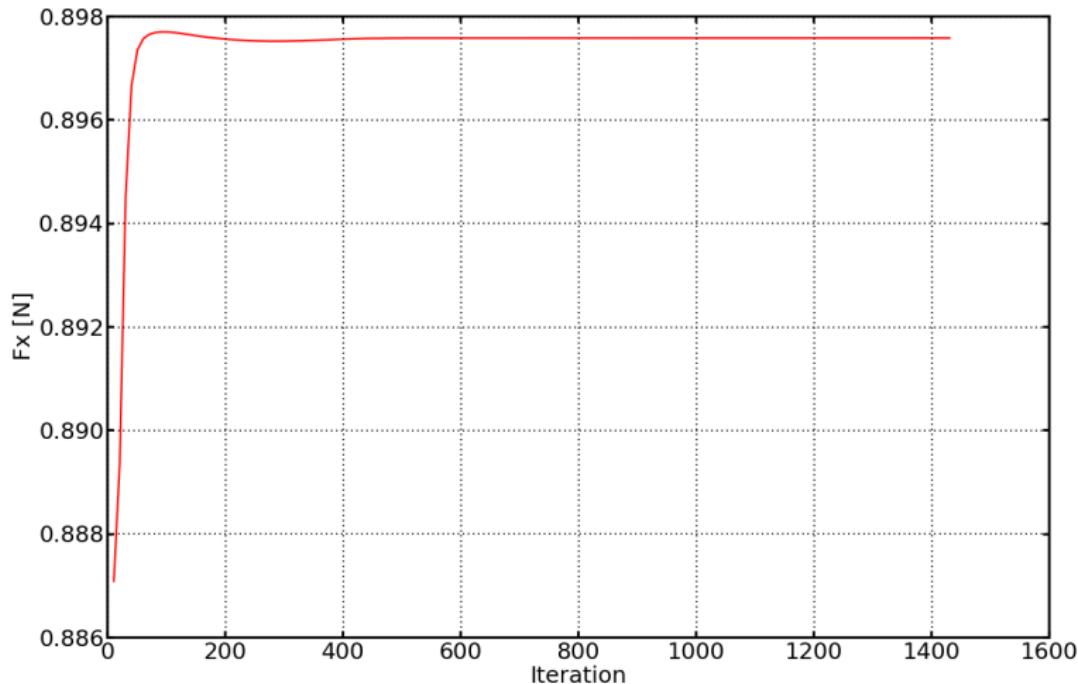
Refined grid - residuals



Original - drag force



Refined grid - drag force



What does this tell us?

- ▶ Many iterations spent on converging the solution
- ▶ The drag force didn't change much - again, wall function is determining the friction resistance
- ▶ Since the initial condition is already pretty converged, relative to the initial state, the same change in residuals does not correspond to the same change in the drag
- ▶ Could have probably stopped the simulation earlier, obtaining the same result but saving some time!

Under-relaxation

- ▶ Copy the fine mesh flat plate case

Under-relaxation

- ▶ Copy the fine mesh flat plate case

Use

```
cp -r ZPG_wallFunction_fine  
ZPG_wallFunction_fine_lessUnderRel
```

Under-relaxation

- ▶ Copy the fine mesh flat plate case
- ▶ In fvSolution, change the relaxation factor for pressure from 0.3 to 0.55

Under-relaxation

- ▶ Copy the fine mesh flat plate case
- ▶ In fvSolution, change the relaxation factor for pressure from 0.3 to 0.55

Edit

```
gedit system/fvSolution
```

Under-relaxation

- ▶ Copy the fine mesh flat plate case
- ▶ In fvSolution, change the relaxation factor for pressure from 0.3 to 0.55

Amend

From this:

```
relaxationFactors
{
    fields
    {
        p    0.3;
    }
    ...
}
```

To This:

```
relaxationFactors
{
    fields
    {
        p    0.55;
    }
    ...
}
```

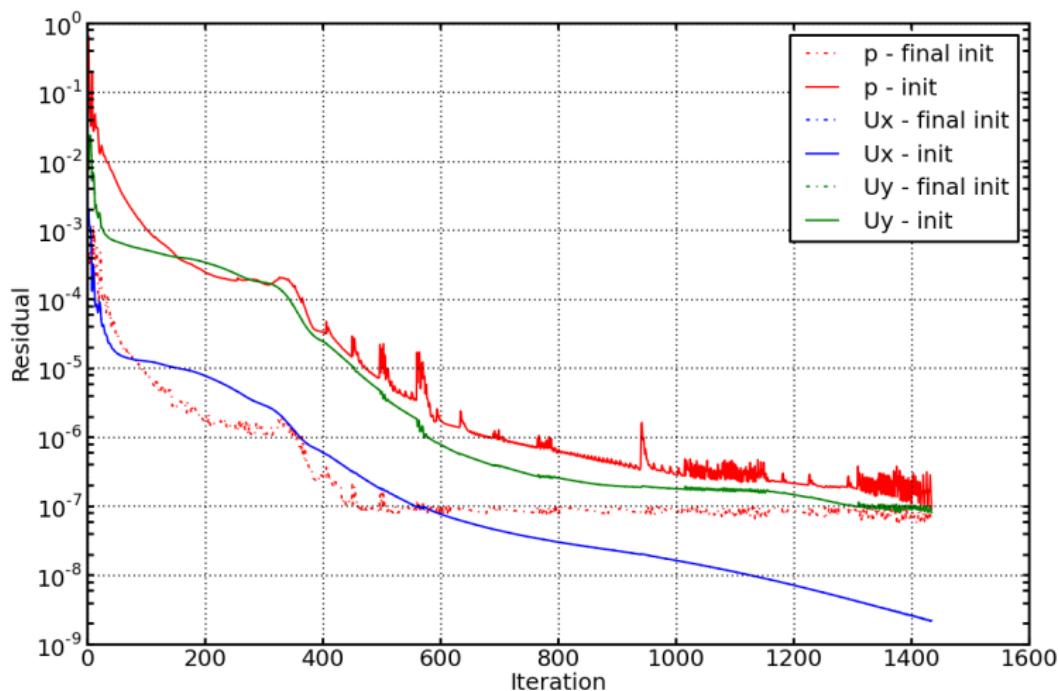
Under-relaxation

- ▶ Copy the fine mesh flat plate case
- ▶ In fvSolution, change the relaxation factor for pressure from 0.3 to 0.55

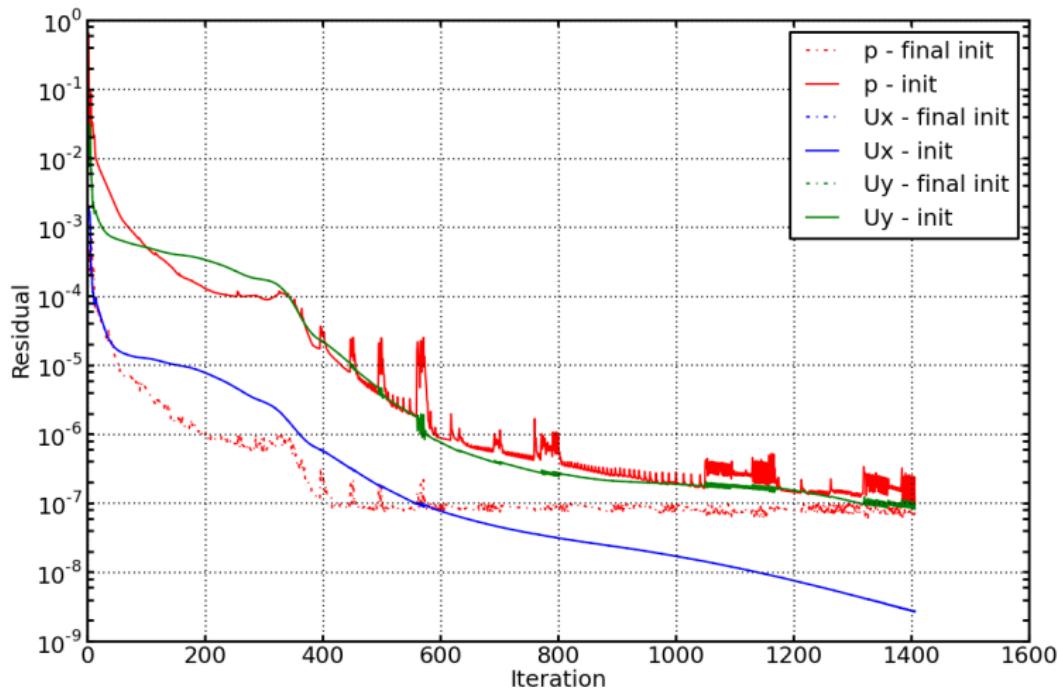
Do

- ▶ Run the simulation
- ▶ Look at the forces and residuals

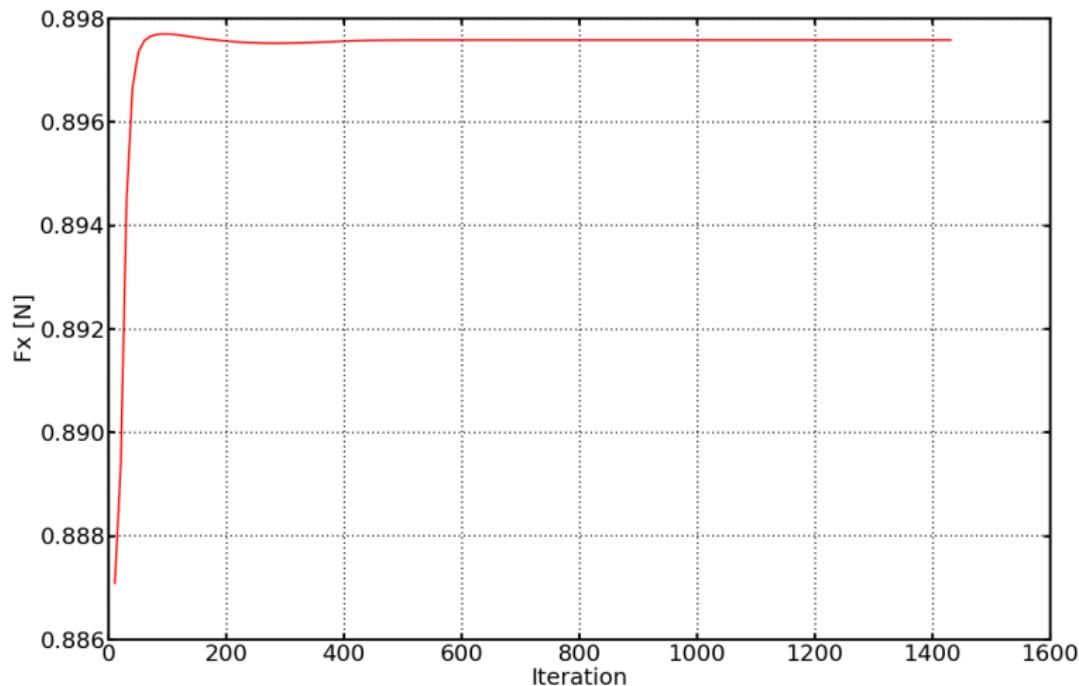
Original (fine mesh) - residuals



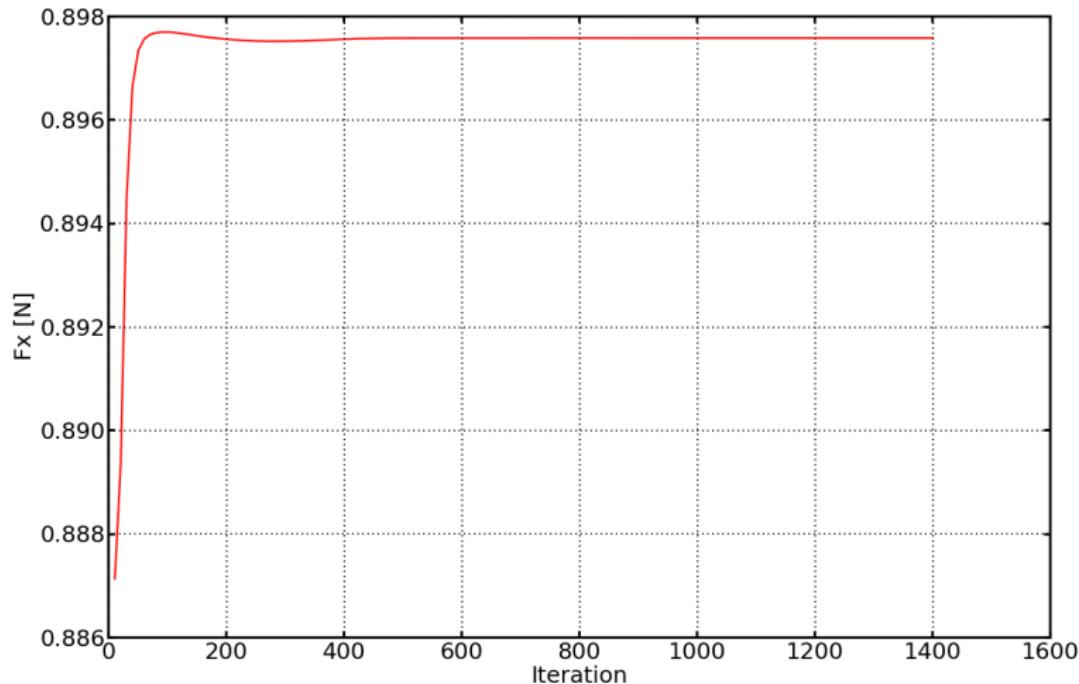
Less under-relaxation - residuals



Original (fine mesh) - drag force



Less under-relaxation - drag force



What does this tell us?

- ▶ We removed the lag in convergence between pressure and wall-normal velocity
- ▶ This means that we are wasting less time looping over already converged U_y equation when p is still changing
- ▶ Very powerful parameter in terms of optimising a solution procedure
- ▶ Cannot push it to far without losing stability
- ▶ Note the more prominent local fluctuations in the pressure equation - they will only get worse the less we under-relax, ultimately causing a crash

High order schemes

- ▶ Again, copy the fine mesh flat plate case

High order schemes

- ▶ Again, copy the fine mesh flat plate case

Use

```
cp -r ZPG_wallFunction_fine  
ZPG_wallFunction_fine_linear
```

High order schemes

- ▶ Again, copy the fine mesh flat plate case
- ▶ In fvSchemes, change the convection scheme from upwind to linear

Edit

gedit system/fvSchemes

High order schemes

- ▶ Again, copy the fine mesh flat plate case
- ▶ In fvSchemes, change the convection scheme from upwind to linear

Amend

- ▶ From this:

```
div(phi,U) bounded Gauss upwind;
```

- ▶ To this:

```
div(phi,U) bounded Gauss linear;
```

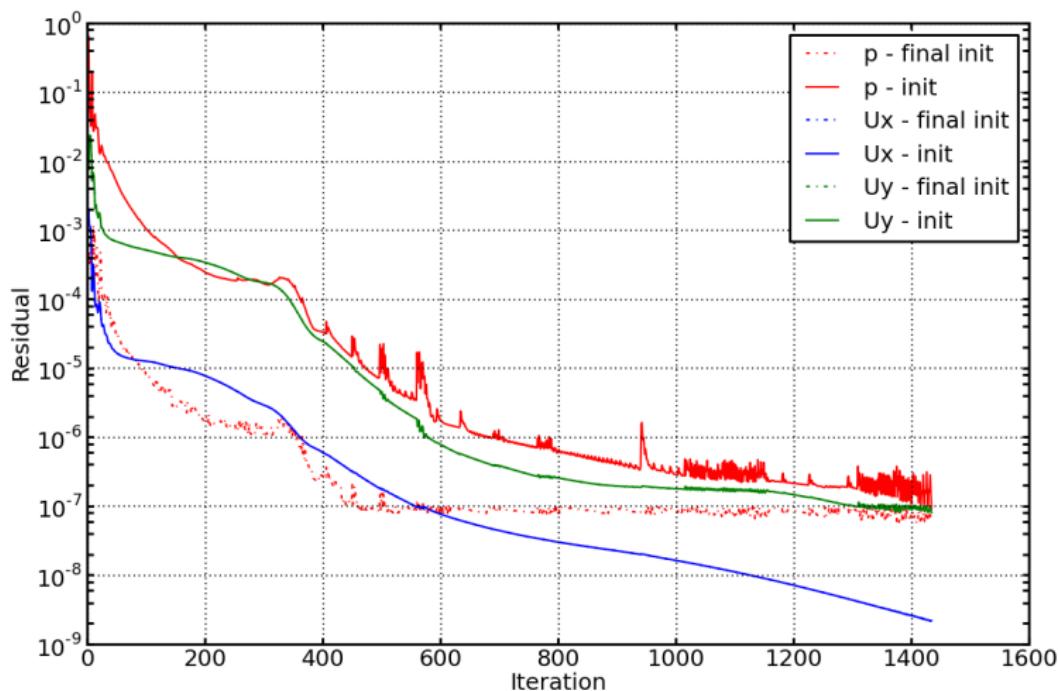
High order schemes

- ▶ Again, copy the fine mesh flat plate case
- ▶ In fvSchemes, change the convection scheme from upwind to linear

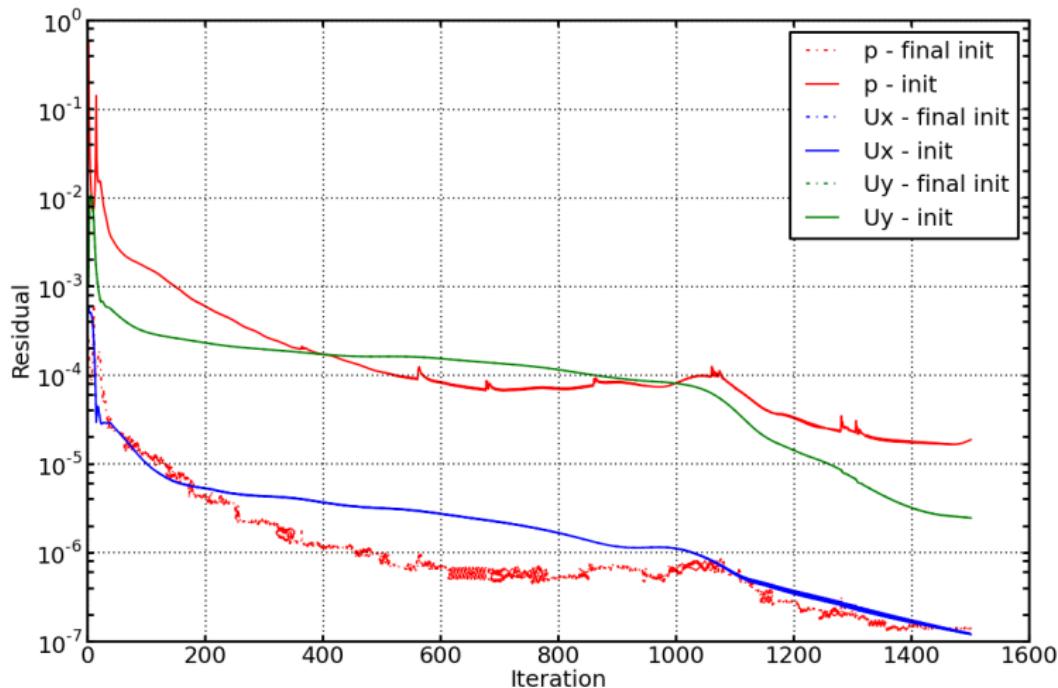
Do

- ▶ Run the simulation
- ▶ Look at the forces and residuals

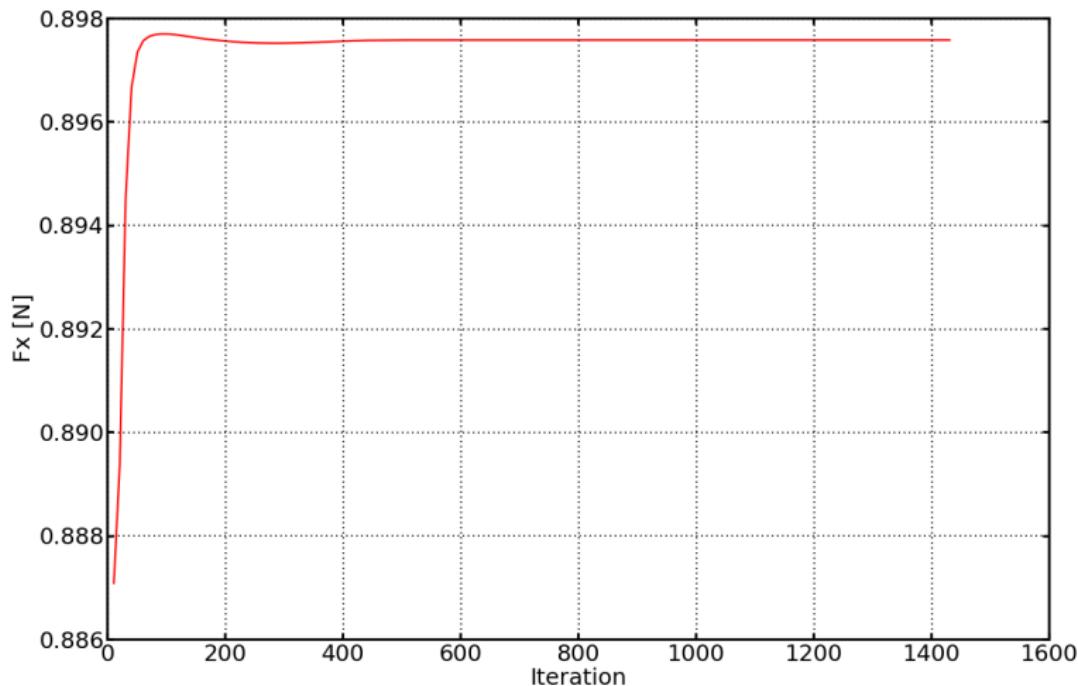
Original (fine mesh) - residuals



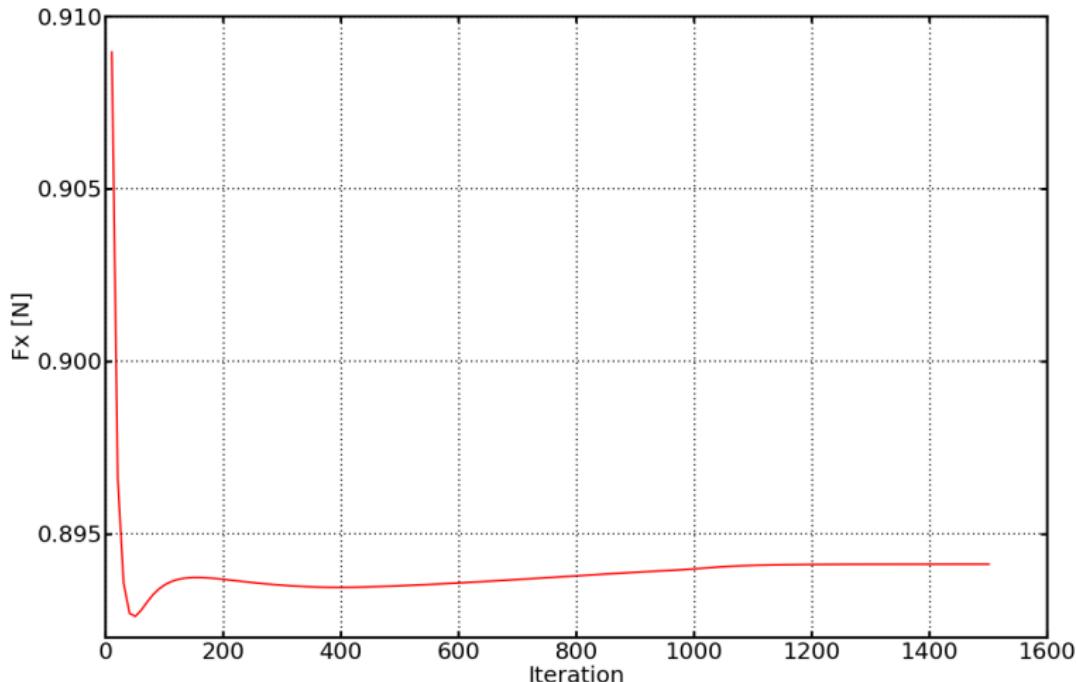
Linear convection scheme - residuals



Original (fine mesh) - drag force



Linear convection scheme - drag force



What does this tell us?

- ▶ Significantly slower convergence, both in terms of iterations and total computational time
- ▶ After the maximum no. iterations the solution is still far from convergence, would need to run for much longer - no guarantee that it will work at all!
- ▶ This shows the difficulties in using pure high order schemes even for very simple problems in finite volume frameworks

Summary

Conclusions

- ▶ We touched upon how CFD actually works, useful to understand the basics in order to troubleshoot problems
- ▶ The use of high order schemes, mesh refinement and under-relaxation on a steady-state problem has been shown to affect the solution
- ▶ The convergence process for an unsteady problem has also been discussed briefly

Further reading

- ▶ Tips and tricks to choosing solvers and schemes in OpenFOAM by Joel Guerrero
[http://www.dicat.unige.it/guerrero/of2014b/
12tipsandtricks.pdf](http://www.dicat.unige.it/guerrero/of2014b/12tipsandtricks.pdf)
- ▶ Implicit vs Explicit time schemes explained
[http://www.cfd-online.com/Forums/main/
5982-implicit-vs-explicit-method.html](http://www.cfd-online.com/Forums/main/5982-implicit-vs-explicit-method.html)
- ▶ Brief introduction to the PISO algorithm
[http://openfoamwiki.net/index.php/OpenFOAM_guide/
The_PISO_algorithm_in_OpenFOAM](http://openfoamwiki.net/index.php/OpenFOAM_guide/The_PISO_algorithm_in_OpenFOAM)
- ▶ Brief introduction to the SIMPLE algorithm
[https://openfoamwiki.net/index.php/OpenFOAM_
guide/The_SIMPLE_algorithm_in_OpenFOAM](https://openfoamwiki.net/index.php/OpenFOAM_guide/The_SIMPLE_algorithm_in_OpenFOAM)

The End